

```
In [1]: import numpy as np
import pandas as pd
import sklearn
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
import os
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
import random
import warnings
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, accuracy_score
from sklearn import metrics
np.random.seed(34)
warnings.filterwarnings('ignore')
```

```
In [2]: index_names = ['ENGINE', 'CYCLE']
setting_names = ['SET1', 'SET2', 'SET3']
sensor_names=[ "INLET_TEMP",
"LPC_OUT_TEMP",
"HPC_OUT_TEMP",
"LPT_OUT_TEMP",
"FAN_IN_PR",
"BYPASS-PR",
"HPC_OUT_PR",
"FAN_RPM",
"CORE_RPM",
"ENGINE_PR",
"HPC_OUT",
"FUEL_RATIO",
"FAN_RPM_CORR",
"CORE_RPM_CORR",
"BYPASS_RATIO",
"FUEL_RATIO_BURNER",
"ENTHALPY_BLEED",
"FAN_SPEED_REQ",
"CONV_FAN_SPEED",
"HP_AIRFLOW",
"LPC_AIRFLOW" ]
col_names = index_names + setting_names + sensor_names
```


```
In [3]: data = pd.read_csv('CMaps\\train_FD001.txt', sep= " ", header=None, index_col=False, na
```

```
In [4]: pd.set_option('display.max_columns', None)
df = data.copy()
df
```

Out[4]:

	ENGINE	CYCLE	SET1	SET2	SET3	INLET_TEMP	LPC_OUT_TEMP	HPC_OUT_TE
0	1	1	-0.0007	-0.0004	100.0	518.67	641.82	158%
1	1	2	0.0019	-0.0003	100.0	518.67	642.15	159%
2	1	3	-0.0043	0.0003	100.0	518.67	642.35	158%
3	1	4	0.0007	0.0000	100.0	518.67	642.35	158%
4	1	5	-0.0019	-0.0002	100.0	518.67	642.37	158%
...
20626	100	196	-0.0004	-0.0003	100.0	518.67	643.49	159%
20627	100	197	-0.0016	-0.0005	100.0	518.67	643.54	160%
20628	100	198	0.0004	0.0000	100.0	518.67	643.42	160%
20629	100	199	-0.0011	0.0003	100.0	518.67	643.23	160%
20630	100	200	-0.0032	-0.0005	100.0	518.67	643.85	160%

20631 rows × 26 columns



```
In [5]: df.shape
```

Out[5]: (20631, 26)

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20631 entries, 0 to 20630
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ENGINE                20631 non-null  int64
1   CYCLE                 20631 non-null  int64
2   SET1                  20631 non-null  float64
3   SET2                  20631 non-null  float64
4   SET3                  20631 non-null  float64
5   INLET_TEMP            20631 non-null  float64
6   LPC_OUT_TEMP          20631 non-null  float64
7   HPC_OUT_TEMP          20631 non-null  float64
8   LPT_OUT_TEMP          20631 non-null  float64
9   FAN_IN_PR             20631 non-null  float64
10  BYPASS-PR             20631 non-null  float64
11  HPC_OUT_PR            20631 non-null  float64
12  FAN_RPM               20631 non-null  float64
13  CORE_RPM              20631 non-null  float64
14  ENGINE_PR             20631 non-null  float64
15  HPC_OUT               20631 non-null  float64
16  FUEL_RATIO            20631 non-null  float64
17  FAN_RPM_CORR          20631 non-null  float64
18  CORE_RPM_CORR         20631 non-null  float64
19  BYPASS_RATIO          20631 non-null  float64
20  FUEL_RATIO_BURNER     20631 non-null  float64
21  ENTHALPY_BLEED        20631 non-null  int64
22  FAN_SPEED_REQ          20631 non-null  int64
23  CONV_FAN_SPEED        20631 non-null  float64
24  HP_AIRFLOW            20631 non-null  float64
25  LPC_AIRFLOW           20631 non-null  float64
dtypes: float64(22), int64(4)
memory usage: 4.1 MB
```

```
In [7]: df.loc[:,['ENGINE','CYCLE']].describe()
```

```
Out[7]:
```

	ENGINE	CYCLE
count	20631.000000	20631.000000
mean	51.506568	108.807862
std	29.227633	68.880990
min	1.000000	1.000000
25%	26.000000	52.000000
50%	52.000000	104.000000
75%	77.000000	156.000000
max	100.000000	362.000000

```
In [8]: df.loc[:,['SET1','SET2','SET3']].describe()
```

Out[8]:

	SET1	SET2	SET3
count	20631.000000	20631.000000	20631.0
mean	-0.000009	0.000002	100.0
std	0.002187	0.000293	0.0
min	-0.008700	-0.000600	100.0
25%	-0.001500	-0.000200	100.0
50%	0.000000	0.000000	100.0
75%	0.001500	0.000300	100.0
max	0.008700	0.000600	100.0

In [9]: `df.loc[:, 'INLET_TEMP': 'LPC_AIRFLOW'].describe()`

Out[9]:

	INLET_TEMP	LPC_OUT_TEMP	HPC_OUT_TEMP	LPT_OUT_TEMP	FAN_IN_PR	BY
count	2.063100e+04	20631.000000	20631.000000	20631.000000	2.063100e+04	2063
mean	5.186700e+02	642.680934	1590.523119	1408.933782	1.462000e+01	2
std	6.537152e-11	0.500053	6.131150	9.000605	3.394700e-12	
min	5.186700e+02	641.210000	1571.040000	1382.250000	1.462000e+01	2
25%	5.186700e+02	642.325000	1586.260000	1402.360000	1.462000e+01	2
50%	5.186700e+02	642.640000	1590.100000	1408.040000	1.462000e+01	2
75%	5.186700e+02	643.000000	1594.380000	1414.555000	1.462000e+01	2
max	5.186700e+02	644.530000	1616.910000	1441.490000	1.462000e+01	2

```
In [10]: unwanted=[]
for i in df.select_dtypes(include=np.number):
    if df[i].nunique() == 1:
        unwanted.append(i)

unwanted
```

```
Out[10]: ['SET3',
          'INLET_TEMP',
          'FAN_IN_PR',
          'ENGINE_PR',
          'FUEL_RATIO_BURNER',
          'FAN_SPEED_REQ',
          'CONV_FAN_SPEED']
```

```
In [11]: df.drop(columns=unwanted, inplace=True)
print(df.shape)
```


df

(20631, 19)

Out[11]:

	ENGINE	CYCLE	SET1	SET2	LPC_OUT_TEMP	HPC_OUT_TEMP	LPT_OUT_TEMP
0	1	1	-0.0007	-0.0004	641.82	1589.70	1400.60
1	1	2	0.0019	-0.0003	642.15	1591.82	1403.14
2	1	3	-0.0043	0.0003	642.35	1587.99	1404.20
3	1	4	0.0007	0.0000	642.35	1582.79	1401.87
4	1	5	-0.0019	-0.0002	642.37	1582.85	1406.22
...
20626	100	196	-0.0004	-0.0003	643.49	1597.98	1428.63
20627	100	197	-0.0016	-0.0005	643.54	1604.50	1433.58
20628	100	198	0.0004	0.0000	643.42	1602.46	1428.18
20629	100	199	-0.0011	0.0003	643.23	1605.26	1426.53
20630	100	200	-0.0032	-0.0005	643.85	1600.38	1432.14

20631 rows × 19 columns



```
In [12]: data_RUL = df.groupby(['ENGINE']).agg({'CYCLE': 'max'})
data_RUL.rename(columns={'CYCLE': 'LIFE'}, inplace=True)
data_RUL.head()
```

Out[12]:


	LIFE
ENGINE	
1	192
2	287
3	179
4	189
5	269

```
In [13]: df_RUL=df.merge(data_RUL,how='left',on=['ENGINE'])
df_RUL['RUL']=df_RUL['LIFE']-df_RUL['CYCLE']
df_RUL.drop(['LIFE'],axis=1,inplace=True)
df_RUL.drop(['CYCLE'],axis=1,inplace=True)
df_RUL
```

Out[13]:

	ENGINE	SET1	SET2	LPC_OUT_TEMP	HPC_OUT_TEMP	LPT_OUT_TEMP	BYPAS F
0	1	-0.0007	-0.0004	641.82	1589.70	1400.60	21.6
1	1	0.0019	-0.0003	642.15	1591.82	1403.14	21.6
2	1	-0.0043	0.0003	642.35	1587.99	1404.20	21.6
3	1	0.0007	0.0000	642.35	1582.79	1401.87	21.6
4	1	-0.0019	-0.0002	642.37	1582.85	1406.22	21.6
...
20626	100	-0.0004	-0.0003	643.49	1597.98	1428.63	21.6
20627	100	-0.0016	-0.0005	643.54	1604.50	1433.58	21.6
20628	100	0.0004	0.0000	643.42	1602.46	1428.18	21.6
20629	100	-0.0011	0.0003	643.23	1605.26	1426.53	21.6
20630	100	-0.0032	-0.0005	643.85	1600.38	1432.14	21.6


20631 rows × 19 columns



```

In [14]: plt.figure(figsize=(15,15))
sns.set_style("whitegrid", {"axes.facecolor": ".0"})
df_cluster2 = df_RUL.corr()
plot_kws={"s": 1}
sns.heatmap(df_RUL.corr(),
            cmap='RdYlBu',
            annot=True,
            linecolor='lightgrey').set_facecolor('white')

```



```
In [15]: corr = df_RUL.corr()

# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr, dtype=bool))

# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(11, 9))


# Generate a custom diverging colormap
cmap = sns.diverging_palette(230, 20, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=1, vmin=-1, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5}, annot=True, fmt=".",
            annot_kws={"size": 8})

# Enhance readability
```

```
plt.xticks(fontsize=10, rotation=45, ha='right')
plt.yticks(fontsize=10)
plt.title('Correlation Heatmap', fontsize=20)

# Show the heatmap
plt.show()
```



```
In [16]: correlations_with_RUL = corr['RUL'].drop('RUL')

# Calculate the absolute values of the correlations
absolute_correlations = correlations_with_RUL.abs()

# Sort the absolute correlations in descending order to identify the strongest corr
sorted_indices = absolute_correlations.sort_values(ascending=False).index

# Retrieve the original correlation values for the sorted indices
sorted_correlations = correlations_with_RUL.loc[sorted_indices]

# Print the sorted correlations with their original signs
print(sorted_correlations)
```



```

HPC_OUT          -0.696228
LPT_OUT_TEMP     -0.678948
FUEL_RATIO       0.671983
HPC_OUT_PR       0.657223
BYPASS_RATIO     -0.642667
LPC_AIRFLOW      0.635662
HP_AIRFLOW       0.629428
LPC_OUT_TEMP     -0.606484
ENTHALPY_BLEED   -0.606154
HPC_OUT_TEMP     -0.584520
FAN_RPM          -0.563968
FAN_RPM_CORR     -0.562569
CORE_RPM         -0.390102
CORE_RPM_CORR    -0.306769
BYPASS-PR        -0.128348
ENGINE           0.078753
SET1             -0.003198
SET2             -0.001948
Name: RUL, dtype: float64

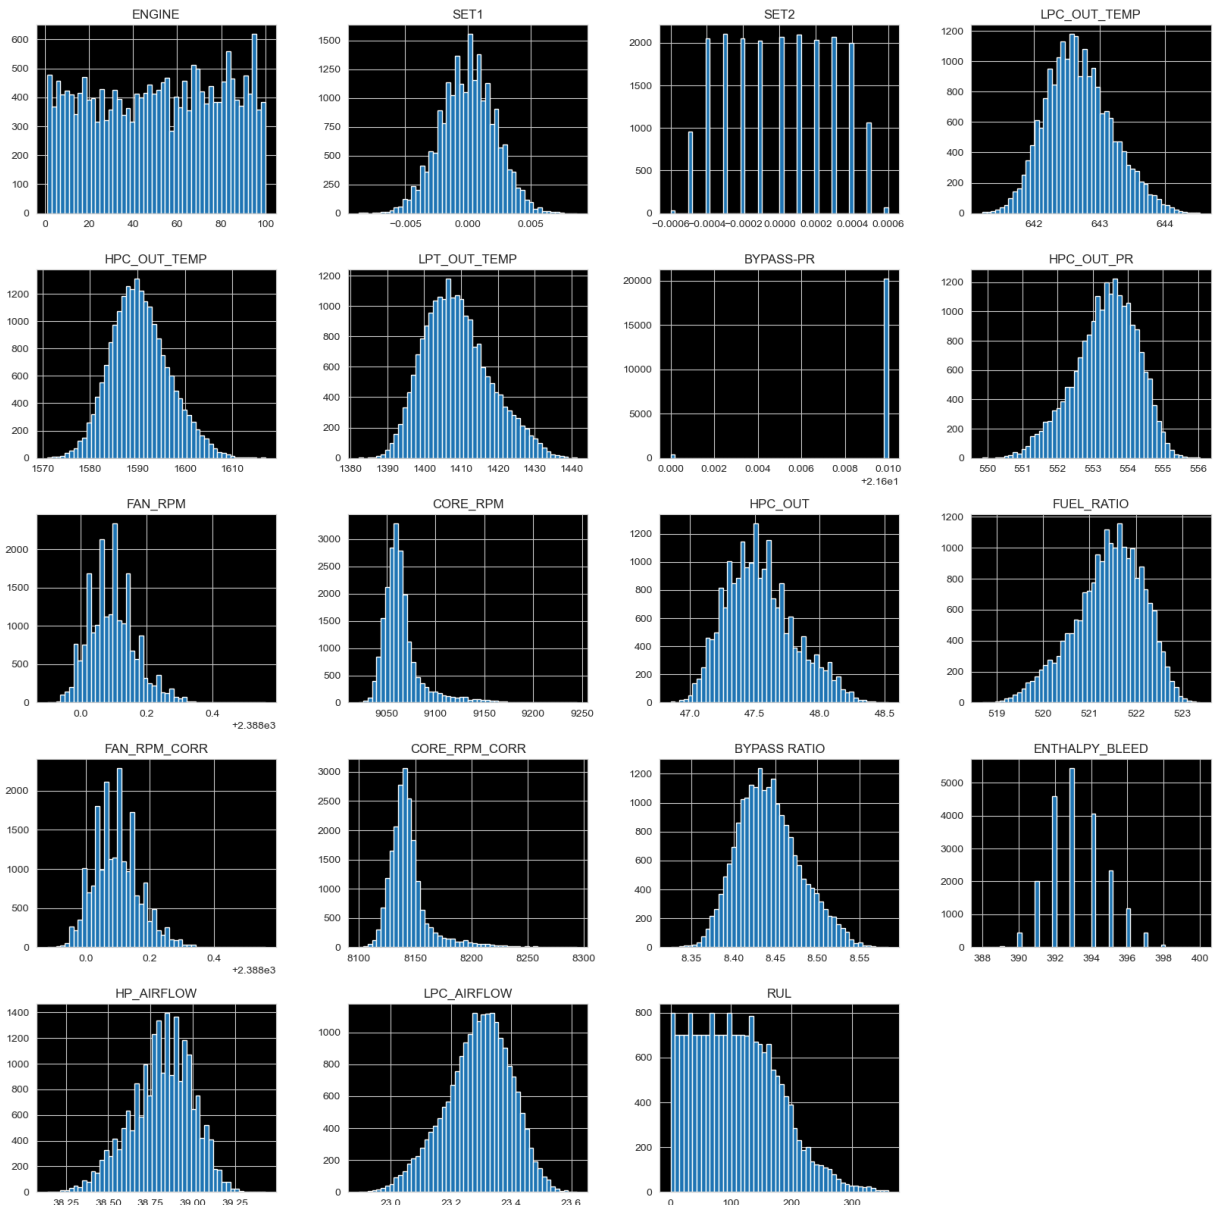
```

```
In [17]: df_RUL.hist(bins=50, figsize=(20, 20))
```

```

Out[17]: array([[<Axes: title={'center': 'ENGINE'}>,
  <Axes: title={'center': 'SET1'}>,
  <Axes: title={'center': 'SET2'}>,
  <Axes: title={'center': 'LPC_OUT_TEMP'}>],
 [<Axes: title={'center': 'HPC_OUT_TEMP'}>,
  <Axes: title={'center': 'LPT_OUT_TEMP'}>,
  <Axes: title={'center': 'BYPASS-PR'}>,
  <Axes: title={'center': 'HPC_OUT_PR'}>],
 [<Axes: title={'center': 'FAN_RPM'}>,
  <Axes: title={'center': 'CORE_RPM'}>,
  <Axes: title={'center': 'HPC_OUT'}>,
  <Axes: title={'center': 'FUEL_RATIO'}>],
 [<Axes: title={'center': 'FAN_RPM_CORR'}>,
  <Axes: title={'center': 'CORE_RPM_CORR'}>,
  <Axes: title={'center': 'BYPASS_RATIO'}>,
  <Axes: title={'center': 'ENTHALPY_BLEED'}>],
 [<Axes: title={'center': 'HP_AIRFLOW'}>,
  <Axes: title={'center': 'LPC_AIRFLOW'}>,
  <Axes: title={'center': 'RUL'}>, <Axes: >]], dtype=object)

```



```
In [18]: labels={1:"DANGER",2:"NOT DANGER YET",3:"SAFE"}

label=[]

#--Transforming rul values to classes :
for i in df_RUL['RUL']:
    if i<=69:
        label.append(1)
    elif i>69 and i<=135:
        label.append(2)
    else:
        label.append(3)
label=np.array(label)

drop_labels = ['ENGINE', 'SET1', 'SET2']
df_train_test=df_RUL.drop(columns=drop_labels).copy()
# Make a copy of the dataset to apply transformations
df_transformed = df_train_test.copy()
X_train, X_test, y_train, y_test=train_test_split(df_train_test,np.array(label), te
```

```
In [19]: #scaler = StandardScaler()
scaler = MinMaxScaler()
#Dropping the target variable
X_train.drop(columns=['RUL'], inplace=True)
X_test.drop(columns=['RUL'], inplace=True)


X_train_scaled=scaler.fit_transform(X_train)
X_test_scaled=scaler.fit_transform(X_test)
```

```
In [20]: X_train_scaled_df = pd.DataFrame(X_train, columns=X_train.columns)

X_train_scaled_df = pd.DataFrame(X_train_scaled, columns=X_train.columns)

X_train_scaled_df.hist(bins=50, figsize=(20, 20))
```

```
Out[20]: array([[<Axes: title={'center': 'LPC_OUT_TEMP'}>,
<Axes: title={'center': 'HPC_OUT_TEMP'}>,
<Axes: title={'center': 'LPT_OUT_TEMP'}>,
<Axes: title={'center': 'BYPASS-PR'}>],
[<Axes: title={'center': 'HPC_OUT_PR'}>,
<Axes: title={'center': 'FAN_RPM'}>,
<Axes: title={'center': 'CORE_RPM'}>,
<Axes: title={'center': 'HPC_OUT'}>]],
[<Axes: title={'center': 'FUEL_RATIO'}>,
<Axes: title={'center': 'FAN_RPM_CORR'}>,
<Axes: title={'center': 'CORE_RPM_CORR'}>,
<Axes: title={'center': 'BYPASS_RATIO'}>],
[<Axes: title={'center': 'ENTHALPY_BLEED'}>,
<Axes: title={'center': 'HP_AIRFLOW'}>,
<Axes: title={'center': 'LPC_AIRFLOW'}>, <Axes: >]], dtype=object)
```



```
In [21]: def mean_absolute_error(y_true, y_pred):
          y_true, y_pred = np.array(y_true), np.array(y_pred)
          return np.mean(np.abs((y_true - y_pred) / (y_true)))
```

```
In [22]: from sklearn.svm import SVC
          classifier_SVC = SVC(kernel = 'linear', random_state = 42)
          classifier=classifier_SVC
          classifier.fit(X_train_scaled,np.array(y_train))
          y_svc_train=classifier.predict(X_train_scaled)

          y_svc_test=classifier.predict(X_test_scaled)

          cm= confusion_matrix(y_test, y_svc_test, labels=classifier.classes_)
          disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=[labels[classifie
          disp.plot()
          plt.grid(False)
          plt.show()
```

```
# Measure the performance
print('SVM-linear')
print("Accuracy score of training %.3f" %metrics.accuracy_score(y_train, y_svc_train))
print("Error rate of training %.3f" %mean_absolute_error(y_train,y_svc_train))
print("Accuracy score of test %.3f" %metrics.accuracy_score(y_test, y_svc_test))
print("Error rate of test %.3f" %mean_absolute_error(y_test,y_svc_test))
print(classification_report(y_test,y_svc_test))
```



SVM-linear

Accuracy score of training 0.686

Error rate of training 0.181

Accuracy score of test 0.660

Error rate of test 0.168

	precision	recall	f1-score	support
1	0.80	0.88	0.84	2800
2	0.49	0.57	0.52	2637
3	0.71	0.53	0.61	2816
accuracy			0.66	8253
macro avg	0.67	0.66	0.66	8253
weighted avg	0.67	0.66	0.66	8253

```
In [23]: from sklearn.ensemble import RandomForestClassifier
classifier_RF=RandomForestClassifier(n_estimators=10)
classifier=classifier_RF
classifier.fit(X_train_scaled,np.array(y_train))
y_svc_train=classifier.predict(X_train_scaled)

y_svc_test=classifier.predict(X_test_scaled)
```

```

cm= confusion_matrix(y_test, y_svc_test, labels=classifier.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=[labels[classifier.classes_]]
disp.plot()
plt.grid(False)
plt.show()

# Measure the performance
print('RandomForestClassifier')
print("Accuracy score of training %.3f" %metrics.accuracy_score(y_train, y_svc_train))
print("Error rate of training %.3f" %mean_absolute_error(y_train,y_svc_train))
print("Accuracy score of test %.3f" %metrics.accuracy_score(y_test, y_svc_test))
print("Error rate of test %.3f" %mean_absolute_error(y_test,y_svc_test))
print(classification_report(y_test,y_svc_test))

```



```

RandomForestClassifier
Accuracy score of training 0.989
Error rate of training 0.006
Accuracy score of test 0.633
Error rate of test 0.187

```


	precision	recall	f1-score	support
1	0.76	0.87	0.81	2800
2	0.47	0.54	0.50	2637
3	0.68	0.48	0.57	2816
accuracy			0.63	8253
macro avg	0.64	0.63	0.63	8253
weighted avg	0.64	0.63	0.63	8253

```
In [24]: from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
classifier=gnb
classifier.fit(X_train_scaled,np.array(y_train))
y_svc_train=classifier.predict(X_train_scaled)

y_svc_test=classifier.predict(X_test_scaled)

cm= confusion_matrix(y_test, y_svc_test, labels=classifier.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=[labels[classifie
disp.plot()
plt.grid(False)
plt.show()

# Measure the performance
print('GaussianNB')
print("Accuracy score of training %.3f" %metrics.accuracy_score(y_train, y_svc_train))
print("Error rate of training %.3f" %mean_absolute_error(y_train,y_svc_train))
print("Accuracy score of test %.3f" %metrics.accuracy_score(y_test, y_svc_test))
print("Error rate of test %.3f" %mean_absolute_error(y_test,y_svc_test))
print(classification_report(y_test,y_svc_test))
```



GaussianNB

Accuracy score of training 0.647

Error rate of training 0.191

Accuracy score of test 0.633

Error rate of test 0.184


	precision	recall	f1-score	support
1	0.73	0.89	0.80	2800
2	0.47	0.52	0.50	2637
3	0.70	0.48	0.57	2816
accuracy			0.63	8253
macro avg	0.64	0.63	0.62	8253
weighted avg	0.64	0.63	0.63	8253

```
In [25]: from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier(n_neighbors=100)
classifier=knn
classifier.fit(X_train_scaled,np.array(y_train))
y_svc_train=classifier.predict(X_train_scaled)

y_svc_test=classifier.predict(X_test_scaled)

cm= confusion_matrix(y_test, y_svc_test, labels=classifier.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=[labels[classifier.classes_]])
disp.plot()
plt.grid(False)
plt.show()

# Measure the performance
print('KNeighborsClassifier')
print("Accuracy score of training %.3f" %metrics.accuracy_score(y_train, y_svc_train))
print("Error rate of training %.3f" %mean_absolute_error(y_train,y_svc_train))
print("Accuracy score of test %.3f" %metrics.accuracy_score(y_test, y_svc_test))
print("Error rate of test %.3f" %mean_absolute_error(y_test,y_svc_test))
print(classification_report(y_test,y_svc_test))
```

KNeighborsClassifier

Accuracy score of training 0.696

Error rate of training 0.187

Accuracy score of test 0.662

Error rate of test 0.174

	precision	recall	f1-score	support
1	0.85	0.84	0.84	2800
2	0.49	0.62	0.55	2637
3	0.70	0.52	0.59	2816
accuracy			0.66	8253
macro avg	0.68	0.66	0.66	8253
weighted avg	0.68	0.66	0.66	8253

```
In [26]: from sklearn.linear_model import LogisticRegression
logistic_regression_model = LogisticRegression(max_iter=1000, random_state=42)
classifier=logistic_regression_model
classifier.fit(X_train_scaled,np.array(y_train))
y_svc_train=classifier.predict(X_train_scaled)

y_svc_test=classifier.predict(X_test_scaled)

cm= confusion_matrix(y_test, y_svc_test, labels=classifier.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=[labels[classifier.classes_]]
disp.plot()
plt.grid(False)
plt.show()
```

```
# Measure the performance
print('LogisticRegression')
print("Accuracy score of training %.3f" %metrics.accuracy_score(y_train, y_svc_train))
print("Error rate of training %.3f" %mean_absolute_error(y_train,y_svc_train))
print("Accuracy score of test %.3f" %metrics.accuracy_score(y_test, y_svc_test))
print("Error rate of test %.3f" %mean_absolute_error(y_test,y_svc_test))
print(classification_report(y_test,y_svc_test))
```



LogisticRegression

Accuracy score of training 0.684

Error rate of training 0.179

Accuracy score of test 0.655

Error rate of test 0.170

	precision	recall	f1-score	support
1	0.77	0.89	0.82	2800
2	0.49	0.55	0.51	2637
3	0.72	0.52	0.61	2816
accuracy			0.66	8253
macro avg	0.66	0.65	0.65	8253
weighted avg	0.66	0.66	0.65	8253

```
In [27]: from sklearn.dummy import DummyClassifier
stratified_clf = DummyClassifier(strategy='stratified', random_state=42)
most_frequent_clf = DummyClassifier(strategy='most_frequent', random_state=42)
uniform_clf = DummyClassifier(strategy='uniform', random_state=42)

baseline_classifiers = {
    'Stratified': stratified_clf,
```

```

    'Most Frequent': most_frequent_clf,
    'Uniform': uniform_clf
}
for name, baseline_clf in baseline_classifiers.items():
    baseline_clf.fit(X_train_scaled, y_train) # Note: Baseline classifiers do not
    y_baseline_pred_test = baseline_clf.predict(X_test_scaled)

    print(f"Evaluation of {name} Classifier:")
    print("Accuracy score of test: %.3f" % accuracy_score(y_test, y_baseline_pred_t
    print(classification_report(y_test, y_baseline_pred_test))
    print("-" * 80)

```

Evaluation of Stratified Classifier:

Accuracy score of test: 0.344

	precision	recall	f1-score	support
1	0.36	0.34	0.35	2800
2	0.34	0.35	0.34	2637
3	0.34	0.34	0.34	2816
accuracy			0.34	8253
macro avg	0.34	0.34	0.34	8253
weighted avg	0.34	0.34	0.34	8253

Evaluation of Most Frequent Classifier:

Accuracy score of test: 0.341

	precision	recall	f1-score	support
1	0.00	0.00	0.00	2800
2	0.00	0.00	0.00	2637
3	0.34	1.00	0.51	2816
accuracy			0.34	8253
macro avg	0.11	0.33	0.17	8253
weighted avg	0.12	0.34	0.17	8253

Evaluation of Uniform Classifier:

Accuracy score of test: 0.329

	precision	recall	f1-score	support
1	0.33	0.33	0.33	2800
2	0.32	0.33	0.32	2637
3	0.33	0.32	0.33	2816
accuracy			0.33	8253
macro avg	0.33	0.33	0.33	8253
weighted avg	0.33	0.33	0.33	8253

```

In [28]: # def update_rolling_mean(data, mask):
#         for x, group in mask.groupby("ENGINE"):
#             for x in X_train.columns:
#                 data.loc[group.index[10:], x+"_rolling"] = data.loc[group.index, x].r

```

```

#         data.loc[group.index[:10], x+"_rolling"] = data.loc[group.index[:10],
# drop_labels = ['BYPASS-PR', 'SET1', 'SET2']
# df_rolling=df_RUL.drop(columns=drop_labels).copy()
# update_rolling_mean(df_rolling, df_rolling)
# df_rolling
def update_rolling_mean(data):
    # Define a function to apply to each group
    def apply_rolling(group):
        for col in group.columns:
            if col != 'ENGINE' and col != 'RUL':
                # Apply rolling mean with a window of 10, min_periods=1 ensures we
                group[col + "_rolling"] = group[col].rolling(window=10, min_periods
        return group

    # Group by 'ENGINE' and apply the rolling function
    data = data.groupby('ENGINE').apply(apply_rolling)
    data.reset_index(inplace=True, drop=True)
    # Optional: drop original columns if you only want to keep the rolling features
    # for col in data.columns:
    #     if not col.endswith('_rolling') and col != 'ENGINE' and col != 'RUL' :
    #         data.drop(col, inplace=True, axis=1)


    return data
df_rolling = df_RUL.copy() # Assuming df_RUL is your initial DataFrame
df_rolling = update_rolling_mean(df_rolling)
df_rolling

```

Out[28]:

	ENGINE	SET1	SET2	LPC_OUT_TEMP	HPC_OUT_TEMP	LPT_OUT_TEMP	BYPAS F
0	1	-0.0007	-0.0004	641.82	1589.70	1400.60	21.0
1	1	0.0019	-0.0003	642.15	1591.82	1403.14	21.0
2	1	-0.0043	0.0003	642.35	1587.99	1404.20	21.0
3	1	0.0007	0.0000	642.35	1582.79	1401.87	21.0
4	1	-0.0019	-0.0002	642.37	1582.85	1406.22	21.0
...
20626	100	-0.0004	-0.0003	643.49	1597.98	1428.63	21.0
20627	100	-0.0016	-0.0005	643.54	1604.50	1433.58	21.0
20628	100	0.0004	0.0000	643.42	1602.46	1428.18	21.0
20629	100	-0.0011	0.0003	643.23	1605.26	1426.53	21.0
20630	100	-0.0032	-0.0005	643.85	1600.38	1432.14	21.0

20631 rows × 36 columns



In [29]:

```

corr = df_rolling.corr()
correlations_with_RUL = corr['RUL'].drop('RUL')

```

```

# Calculate the absolute values of the correlations
absolute_correlations = correlations_with_RUL.abs()

# Sort the absolute correlations in descending order to identify the strongest corr
sorted_indices = absolute_correlations.sort_values(ascending=False).index

# Retrieve the original correlation values for the sorted indices
sorted_correlations = correlations_with_RUL.loc[sorted_indices]
features_to_drop = absolute_correlations[absolute_correlations < 0.1].index
# Print the sorted correlations with their original signs
print(sorted_correlations)
features_to_drop

```

```

LPT_OUT_TEMP_rolling    -0.732868
LPC_AIRFLOW_rolling     0.729701
HPC_OUT_rolling         -0.728727
ENTHALPY_BLEED_rolling  -0.728268
BYPASS_RATIO_rolling    -0.727228
HPC_OUT_TEMP_rolling    -0.726007
HP_AIRFLOW_rolling      0.723255
LPC_OUT_TEMP_rolling    -0.721540
HPC_OUT_PR_rolling      0.711189
FUEL_RATIO_rolling      0.710306
HPC_OUT                 -0.696228
LPT_OUT_TEMP            -0.678948
FUEL_RATIO              0.671983
HPC_OUT_PR              0.657223
BYPASS_RATIO            -0.642667
LPC_AIRFLOW             0.635662
HP_AIRFLOW              0.629428
LPC_OUT_TEMP            -0.606484
ENTHALPY_BLEED          -0.606154
FAN_RPM_rolling         -0.592877
FAN_RPM_CORR_rolling    -0.591497
HPC_OUT_TEMP            -0.584520
FAN_RPM                 -0.563968
FAN_RPM_CORR            -0.562569
CORE_RPM_rolling        -0.393039
CORE_RPM                -0.390102
CORE_RPM_CORR           -0.306769
CORE_RPM_CORR_rolling   -0.305386
BYPASS-PR_rolling       -0.301141
BYPASS-PR               -0.128348
ENGINE                  0.078753
SET2_rolling            -0.016361
SET1_rolling            -0.012197
SET1                    -0.003198
SET2                    -0.001948
Name: RUL, dtype: float64

```

```
Out[29]: Index(['ENGINE', 'SET1', 'SET2', 'SET1_rolling', 'SET2_rolling'], dtype='object')
```


```
In [30]: df_rolling.drop(columns=['SET1', 'SET2', 'SET1_rolling', 'SET2_rolling', 'BYPASS-PR',
X_train, X_test, y_train, y_test=train_test_split(df_rolling.drop(columns=['ENGINE'
```

```
df_rolling.iloc[-1,-14:]=df_rolling.iloc[-2,-14:]
df_rolling
```

Out[30]:

	ENGINE	LPC_OUT_TEMP	HPC_OUT_TEMP	LPT_OUT_TEMP	HPC_OUT_PR	FAN_RPM
0	1	641.82	1589.70	1400.60	554.36	2388.06
1	1	642.15	1591.82	1403.14	553.75	2388.04
2	1	642.35	1587.99	1404.20	554.26	2388.08
3	1	642.35	1582.79	1401.87	554.45	2388.11
4	1	642.37	1582.85	1406.22	554.00	2388.06
...
20626	100	643.49	1597.98	1428.63	551.43	2388.19
20627	100	643.54	1604.50	1433.58	550.86	2388.23
20628	100	643.42	1602.46	1428.18	550.94	2388.24
20629	100	643.23	1605.26	1426.53	550.68	2388.25
20630	100	643.85	1600.38	1432.14	550.79	2388.26

20631 rows × 30 columns




```
In [31]: X_train_scaled=scaler.fit_transform(X_train)
X_test_scaled=scaler.fit_transform(X_test)
```

```
In [32]: classifier_SVC = SVC(kernel = 'linear',random_state = 42)
classifier=classifier_SVC
classifier.fit(X_train_scaled,np.array(y_train))
y_svc_train=classifier.predict(X_train_scaled)

y_svc_test=classifier.predict(X_test_scaled)

cm= confusion_matrix(y_test, y_svc_test, labels=classifier.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=[labels[classifier.classes_]
disp.plot()
plt.grid(False)
plt.show()

# Measure the performance
print('SVM-linear')
print("Accuracy score of training %.3f" %metrics.accuracy_score(y_train, y_svc_train))
print("Error rate of training %.3f" %mean_absolute_error(y_train,y_svc_train))
print("Accuracy score of test %.3f" %metrics.accuracy_score(y_test, y_svc_test))
print("Error rate of test %.3f" %mean_absolute_error(y_test,y_svc_test))
print(classification_report(y_test,y_svc_test))
```



SVM-linear

Accuracy score of training 0.691

Error rate of training 0.178

Accuracy score of test 0.681

Error rate of test 0.204

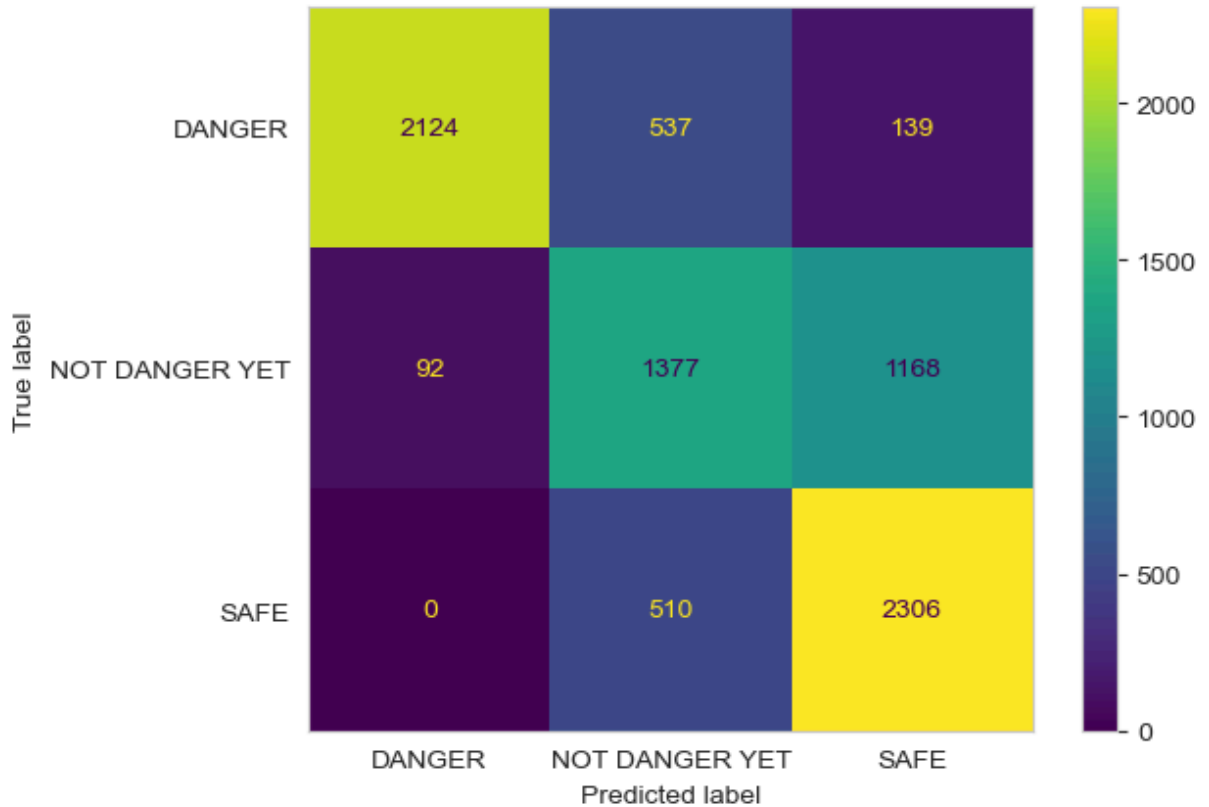
	precision	recall	f1-score	support
1	0.96	0.73	0.83	2800
2	0.52	0.53	0.52	2637
3	0.64	0.77	0.70	2816
accuracy			0.68	8253
macro avg	0.70	0.68	0.68	8253
weighted avg	0.71	0.68	0.69	8253

```
In [33]: from sklearn.svm import SVC
classifier_SVC = SVC(kernel = 'poly', random_state = 42)
classifier=classifier_SVC
classifier.fit(X_train_scaled,np.array(y_train))
y_svc_train=classifier.predict(X_train_scaled)

y_svc_test=classifier.predict(X_test_scaled)

cm= confusion_matrix(y_test, y_svc_test, labels=classifier.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=[labels[classifier.classes_]]
disp.plot()
plt.grid(False)
plt.show()
```

```
# Measure the performance
print('SVM-Ploy')
print("Accuracy score of training %.3f" %metrics.accuracy_score(y_train, y_svc_train))
print("Error rate of training %.3f" %mean_absolute_error(y_train,y_svc_train))
print("Accuracy score of test %.3f" %metrics.accuracy_score(y_test, y_svc_test))
print("Error rate of test %.3f" %mean_absolute_error(y_test,y_svc_test))
print(classification_report(y_test,y_svc_test))
```



SVM-Ploy

Accuracy score of training 0.725

Error rate of training 0.172

Accuracy score of test 0.704

Error rate of test 0.196

	precision	recall	f1-score	support
1	0.96	0.76	0.85	2800
2	0.57	0.52	0.54	2637
3	0.64	0.82	0.72	2816
accuracy			0.70	8253
macro avg	0.72	0.70	0.70	8253
weighted avg	0.72	0.70	0.71	8253

```
In [34]: logistic_regression_model = LogisticRegression(max_iter=1000, random_state=42)
classifier=logistic_regression_model
classifier.fit(X_train_scaled,np.array(y_train))
y_svc_train=classifier.predict(X_train_scaled)

y_svc_test=classifier.predict(X_test_scaled)
```




```

cm= confusion_matrix(y_test, y_svc_test, labels=classifier.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=[labels[classifier.classes_]]
disp.plot()
plt.grid(False)
plt.show()

# Measure the performance
print('LogisticRegression')
print("Accuracy score of training %.3f" %metrics.accuracy_score(y_train, y_svc_train))
print("Error rate of training %.3f" %mean_absolute_error(y_train,y_svc_train))
print("Accuracy score of test %.3f" %metrics.accuracy_score(y_test, y_svc_test))
print("Error rate of test %.3f" %mean_absolute_error(y_test,y_svc_test))
print(classification_report(y_test,y_svc_test))

```



```

LogisticRegression
Accuracy score of training 0.690
Error rate of training 0.174
Accuracy score of test 0.692
Error rate of test 0.190

```

	precision	recall	f1-score	support
1	0.93	0.78	0.85	2800
2	0.54	0.50	0.52	2637
3	0.64	0.78	0.71	2816
accuracy			0.69	8253
macro avg	0.70	0.69	0.69	8253
weighted avg	0.71	0.69	0.69	8253

```

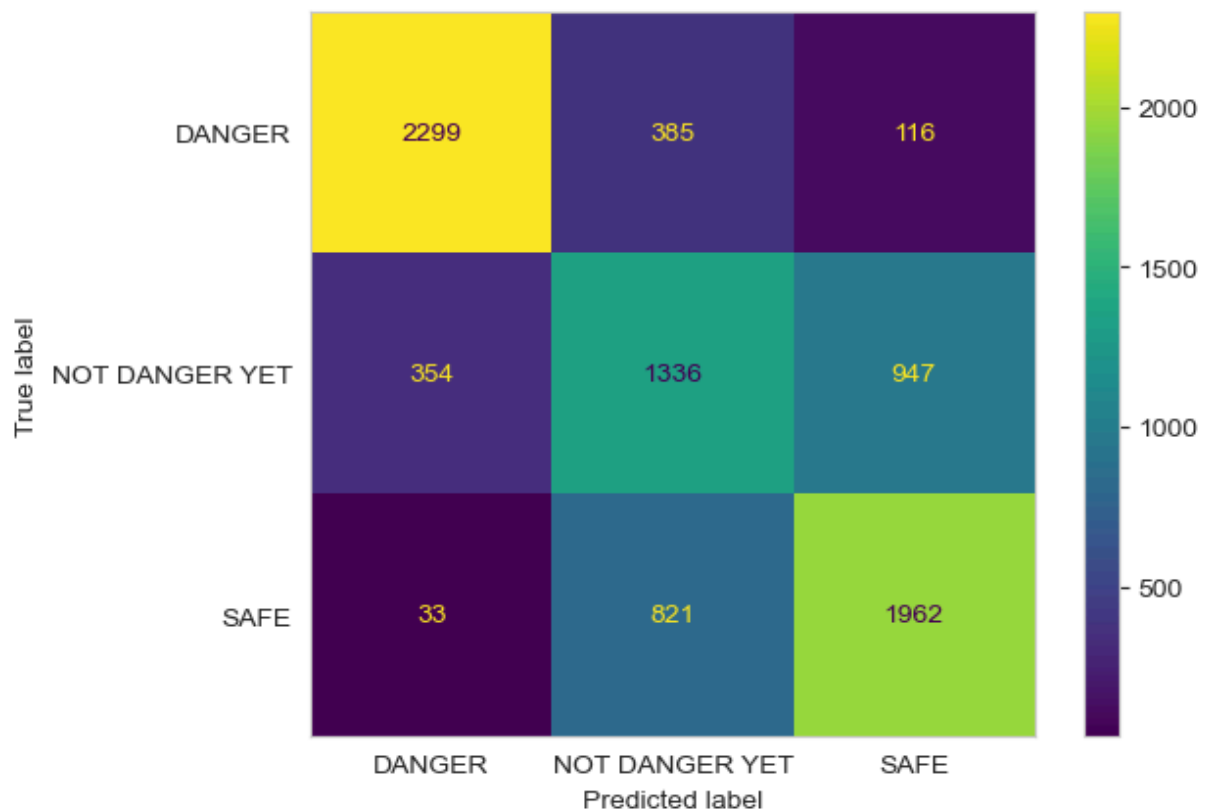
In [35]: classifier_RF=RandomForestClassifier(n_estimators=10)
classifier=classifier_RF
classifier.fit(X_train_scaled,np.array(y_train))
y_svc_train=classifier.predict(X_train_scaled)

y_svc_test=classifier.predict(X_test_scaled)

cm= confusion_matrix(y_test, y_svc_test, labels=classifier.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=[labels[classifie
disp.plot()
plt.grid(False)
plt.show()

# Measure the performance
print('RandomForestClassifier')
print("Accuracy score of training %.3f" %metrics.accuracy_score(y_train, y_svc_train))
print("Error rate of training %.3f" %mean_absolute_error(y_train,y_svc_train))
print("Accuracy score of test %.3f" %metrics.accuracy_score(y_test, y_svc_test))
print("Error rate of test %.3f" %mean_absolute_error(y_test,y_svc_test))
print(classification_report(y_test,y_svc_test))

```



RandomForestClassifier

Accuracy score of training 0.993

Error rate of training 0.004

Accuracy score of test 0.678

Error rate of test 0.189

	precision	recall	f1-score	support
1	0.86	0.82	0.84	2800
2	0.53	0.51	0.52	2637
3	0.65	0.70	0.67	2816
accuracy			0.68	8253
macro avg	0.68	0.67	0.68	8253
weighted avg	0.68	0.68	0.68	8253

```
In [36]: gnb = GaussianNB()
classifier=gnb
classifier.fit(X_train_scaled,np.array(y_train))
y_svc_train=classifier.predict(X_train_scaled)

y_svc_test=classifier.predict(X_test_scaled)

cm= confusion_matrix(y_test, y_svc_test, labels=classifier.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=[labels[classifier.classes_])
disp.plot()
plt.grid(False)
plt.show()

# Measure the performance
print('GaussianNB')
print("Accuracy score of training %.3f" %metrics.accuracy_score(y_train, y_svc_train))
print("Error rate of training %.3f" %mean_absolute_error(y_train,y_svc_train))
print("Accuracy score of test %.3f" %metrics.accuracy_score(y_test, y_svc_test))
print("Error rate of test %.3f" %mean_absolute_error(y_test,y_svc_test))
print(classification_report(y_test,y_svc_test))
```



GaussianNB

Accuracy score of training 0.638

Error rate of training 0.200

Accuracy score of test 0.645

Error rate of test 0.206

	precision	recall	f1-score	support
1	0.89	0.77	0.83	2800
2	0.46	0.47	0.47	2637
3	0.62	0.68	0.65	2816
accuracy			0.65	8253
macro avg	0.66	0.64	0.65	8253
weighted avg	0.66	0.65	0.65	8253

```
In [37]: knn=KNeighborsClassifier(n_neighbors=100)
classifier=knn
classifier.fit(X_train_scaled,np.array(y_train))
y_svc_train=classifier.predict(X_train_scaled)

y_svc_test=classifier.predict(X_test_scaled)

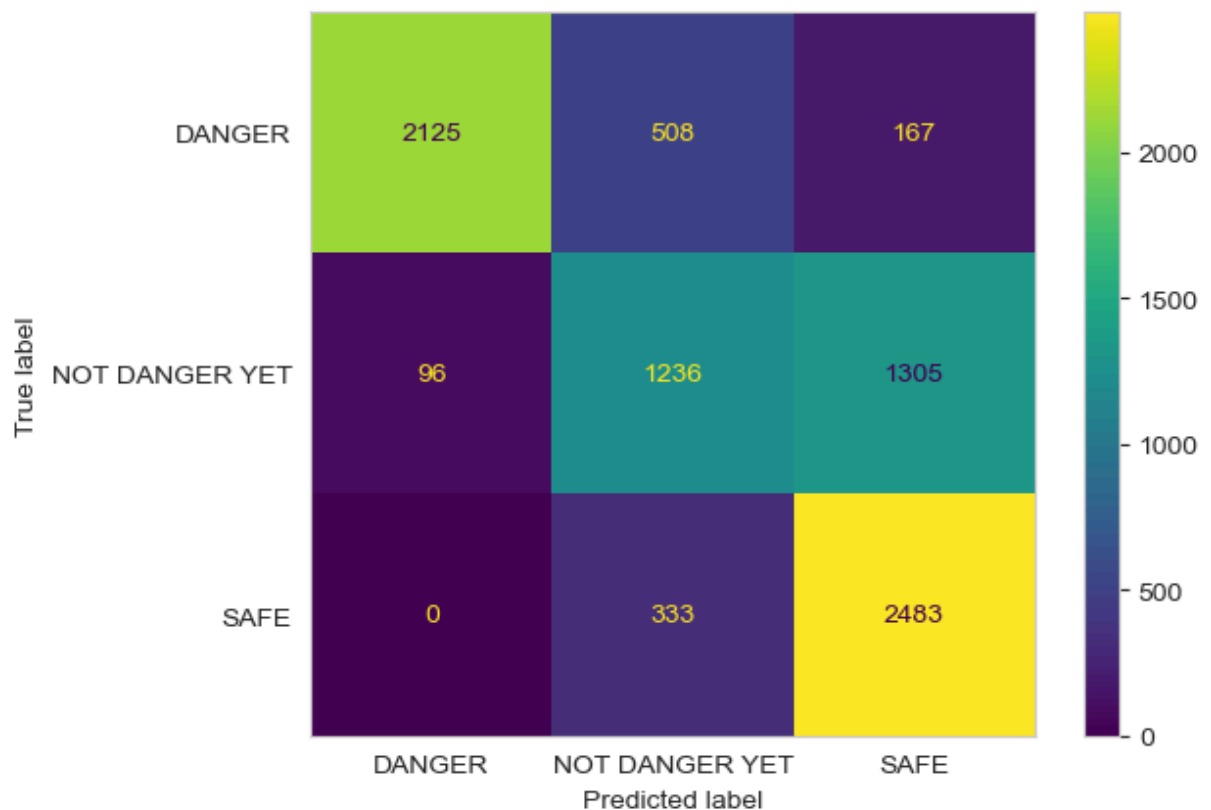
cm= confusion_matrix(y_test, y_svc_test, labels=classifier.classes_)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=[labels[classifier.classes_[0]],
labels[classifier.classes_[1]], labels[classifier.classes_[2]]])
disp.plot()
plt.grid(False)
plt.show()

# Measure the performance
```

```

print('KNeighborsClassifier')
print("Accuracy score of training %.3f" %metrics.accuracy_score(y_train, y_svc_train))
print("Error rate of training %.3f" %mean_absolute_error(y_train,y_svc_train))
print("Accuracy score of test %.3f" %metrics.accuracy_score(y_test, y_svc_test))
print("Error rate of test %.3f" %mean_absolute_error(y_test,y_svc_test))
print(classification_report(y_test,y_svc_test))

```



KNeighborsClassifier

Accuracy score of training 0.723

Error rate of training 0.172

Accuracy score of test 0.708

Error rate of test 0.200

	precision	recall	f1-score	support
1	0.96	0.76	0.85	2800
2	0.60	0.47	0.52	2637
3	0.63	0.88	0.73	2816
accuracy			0.71	8253
macro avg	0.73	0.70	0.70	8253
weighted avg	0.73	0.71	0.70	8253